

# Introduction to SQL

Select-From-Where Statements

Multirelation Queries

Subqueries

# Why SQL?

- SQL is a very-high-level language.
  - Say “what to do” rather than “how to do it.”
  - Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system figures out “best” way to execute query.
  - Called “query optimization.”

# Select-From-Where Statements

**SELECT** desired attributes

**FROM** one or more tables

**WHERE** condition about tuples of  
the tables

# Our Running Example

- All our SQL queries will be based on the following database schema.

- Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)

Frequents(drinker, bar)

# Example

- Using **Beers(name, manf)**, what beers are made by Anheuser-Busch?

```
SELECT name
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch';
```

# Result of Query

name
Bud
Bud Lite
Michelob
...

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.

# Meaning of Single-Relation Query

- Begin with the relation in the FROM clause.
- Apply the selection indicated by the WHERE clause.
- Apply the extended projection indicated by the SELECT clause.

# Operational Semantics

name	manf
Bud	Anheuser-Busch

Tuple-variable  $t$   
loops over all  
tuples

Include  $t.name$   
in the result, if so

Check if  
Anheuser-Busch



# Operational Semantics --- General

- Think of a *tuple variable* visiting each tuple of the relation mentioned in FROM.
- Check if the “current” tuple satisfies the WHERE clause.
- If so, compute the attributes or expressions of the SELECT clause using the components of this tuple.

## \* In SELECT clauses

- When there is one relation in the FROM clause, \* in the SELECT clause stands for “all attributes of this relation.”
- **Example:** Using **Beers(name, manf):**

```
SELECT *
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch';
```

# Result of Query:

name	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
...	...

Now, the result has each of the attributes of Beers.

# Renaming Attributes

- If you want the result to have different attribute names, use "AS <new name>" to rename an attribute.
- **Example:** Using **Beers(name, manf):**

```
SELECT name AS beer, manf
FROM Beers
WHERE manf = 'Anheuser-Busch'
```

# Result of Query:

beer	manf
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
...	...

# Expressions in SELECT Clauses

- Any expression that makes sense can appear as an element of a SELECT clause.
- **Example:** Using `Sells(bar, beer, price)`:

```
SELECT bar, beer,  
       price*114 AS priceInYen  
FROM Sells;
```

# Result of Query

bar	beer	priceInYen	
Joe's	Bud	285	
Sue's	Miller	342	
...	...	...	

# Example: Constants as Expressions

- Using Likes(drinker, beer):

```
SELECT drinker,  
       'likes Bud' AS whoLikesBud  
FROM Likes  
WHERE beer = 'Bud';
```



# Result of Query

drinker who Likes Bud		
Sally	likes	Bud
Fred	likes	Bud
...	...	

# Example: Information Integration

- We often build “data warehouses” from the data at many “sources.”
- Suppose each bar has its own relation `Menu(beer, price)` .
- To contribute to `Sells(bar, beer, price)` we need to query each bar and insert the name of the bar.

# Information Integration --- (2)

- For instance, at Joe's Bar we can issue the query:

```
SELECT 'Joe''s Bar', beer, price  
FROM Menu;
```

# Complex Conditions in WHERE Clause

- Boolean operators AND, OR, NOT.
- Comparisons =, <>, <, >, <=, >=.
- And many other operators that produce boolean-valued results.

# Example: Complex Condition

- Using `Sells(bar, beer, price)`, find the price Joe's Bar charges for Bud:

```
SELECT price
FROM Sells
WHERE bar = 'Joe''s Bar' AND
       beer = 'Bud';
```

# Patterns

- A condition can compare a string to a pattern by:
  - <Attribute> LIKE <pattern> or  
<Attribute> NOT LIKE <pattern>
- *Pattern* is a quoted string with % = "any string"; \_ = "any character."

## Example: LIKE

- Using Drinkers(name, addr, phone) find the drinkers with exchange 555:

```
SELECT name
FROM Drinkers
WHERE phone LIKE '%555-__ __ __';
```

# NULL Values

- Tuples in SQL relations can have NULL as a value for one or more components.
- Meaning depends on context. Two common cases:
  - *Missing value* : e.g., we know Joe's Bar has some address, but we don't know what it is.
  - *Inapplicable* : e.g., the value of attribute *spouse* for an unmarried person.



# Comparing NULL's to Values

- The logic of conditions in SQL is really 3-valued logic: TRUE, FALSE, UNKNOWN.
- Comparing any value (including NULL itself) with NULL yields UNKNOWN.
- A tuple is in a query answer iff the WHERE clause is TRUE (not FALSE or UNKNOWN).

# Three-Valued Logic

- To understand how AND, OR, and NOT work in 3-valued logic, think of TRUE = 1, FALSE = 0, and UNKNOWN =  $\frac{1}{2}$ .
- AND = MIN; OR = MAX, NOT( $x$ ) =  $1-x$ .
- Example:

$$\begin{aligned}\text{TRUE AND (FALSE OR NOT(UNKNOWN))} &= \\ \text{MIN}(1, \text{MAX}(0, (1 - \frac{1}{2}))) &= \\ \text{MIN}(1, \text{MAX}(0, \frac{1}{2})) &= \text{MIN}(1, \frac{1}{2}) = \frac{1}{2}.\end{aligned}$$

# Surprising Example

- From the following Sells relation:

bar	beer	price	
Joe's Bar	Bud	NULL	

SELECT bar

FROM Sells

WHERE price < 2.00 OR price >= 2.00;

← UNKNOWN →      ← UNKNOWN →

← UNKNOWN →

# Reason: 2-Valued Laws $\neq$ 3-Valued Laws

- Some common laws, like commutativity of AND, hold in 3-valued logic.
- But not others, e.g., the *law of the excluded middle* :  $p \text{ OR NOT } p = \text{TRUE}$ .
  - When  $p = \text{UNKNOWN}$ , the left side is  $\text{MAX}( \frac{1}{2}, (1 - \frac{1}{2}) ) = \frac{1}{2} \neq 1$ .

# Multirelation Queries

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by "<relation>.<attribute>".

# Example: Joining Two Relations

- Using relations `Likes(drinker, beer)` and `Frequents(drinker, bar)`, find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer
FROM Likes, Frequents
WHERE bar = 'Joe's Bar' AND
      Frequents.drinker =
      Likes.drinker;
```

# Formal Semantics


- Almost the same as for single-relation queries:
  1. Start with the product of all the relations in the FROM clause.
  2. Apply the selection condition from the WHERE clause.
  3. Project onto the list of attributes and expressions in the SELECT clause.

# Operational Semantics

- Imagine one tuple-variable for each relation in the FROM clause.
  - These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the WHERE clause, send these tuples to the SELECT clause.



# Example



# Explicit Tuple-Variables

- Sometimes, a query needs to use two copies of the same relation.
- Distinguish copies by following the relation name by the name of a tuple-variable, in the FROM clause.
- It's always an option to rename relations this way, even when not essential.

# Example: Self-Join

- From **Beers(name, manf)**, find all pairs of beers by the same manufacturer.
  - Do not produce pairs like (Bud, Bud).
  - Produce pairs in alphabetic order, e.g. (Bud, Miller), not (Miller, Bud).

```
SELECT b1.name, b2.name
FROM Beers b1, Beers b2
WHERE b1.manf = b2.manf AND
      b1.name < b2.name;
```

# Subqueries

- A parenthesized SELECT-FROM-WHERE statement (*subquery*) can be used as a value in a number of places, including FROM and WHERE clauses.
- **Example:** in place of a relation in the FROM clause, we can use a subquery and then query its result.
  - Must use a tuple-variable to name tuples of the result.

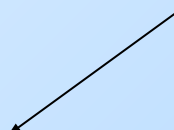
# Example: Subquery in FROM

- Find the beers liked by at least one person who frequents Joe's Bar.

SELECT beer

FROM Likes, (SELECT drinker  
FROM Frequents  
WHERE bar = 'Joe's Bar') JD

Drinkers who  
frequent Joe's Bar



WHERE Likes.drinker = JD.drinker;

# Subqueries That Return One Tuple

- If a subquery is guaranteed to produce one tuple, then the subquery can be used as a value.
  - Usually, the tuple has one component.
  - A run-time error occurs if there is no tuple or more than one tuple.

# Example: Single-Tuple Subquery

- Using `Sells(bar, beer, price)`, find the bars that serve Miller for the same price Joe charges for Bud.
- Two queries would surely work:
  1. Find the price Joe charges for Bud.
  2. Find the bars that serve Miller at that price.

# Query + Subquery Solution

SELECT bar

FROM Sells

WHERE beer = 'Miller' AND

price = (SELECT price

FROM Sells

WHERE bar = 'Joe's Bar'

AND beer = 'Bud');

The price at  
which Joe  
sells Bud



# The IN Operator

- $\langle \text{tuple} \rangle \text{ IN } (\langle \text{subquery} \rangle)$  is true if and only if the tuple is a member of the relation produced by the subquery.
  - Opposite:  $\langle \text{tuple} \rangle \text{ NOT IN } (\langle \text{subquery} \rangle)$ .
- IN-expressions can appear in WHERE clauses.

## Example: IN

- Using **Beers(name, manf)** and **Likes(drinker, beer)**, find the name and manufacturer of each beer that Fred likes.

SELECT \*

FROM Beers

WHERE name IN (SELECT beer

FROM Likes

WHERE drinker = 'Fred');

The set of  
beers Fred  
likes

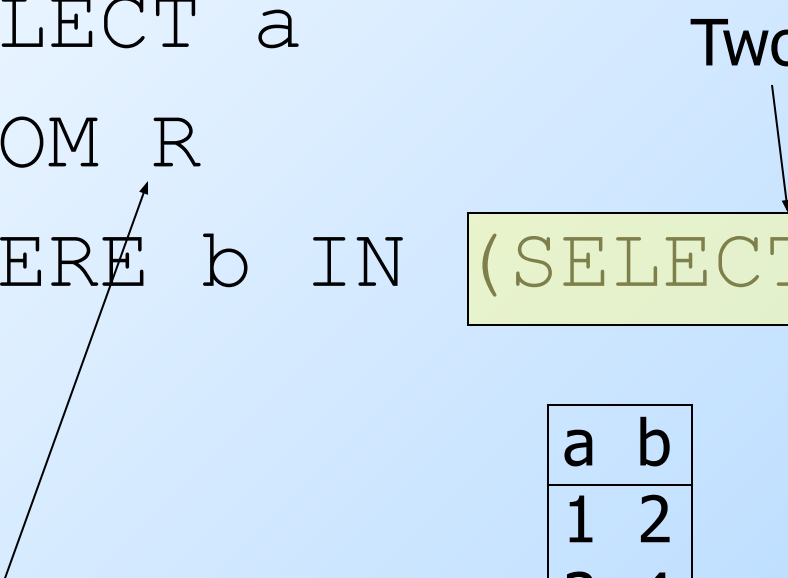
# Remember These From Lecture #1?

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```

# IN is a Predicate About R's Tuples

```
SELECT a
FROM R
WHERE b IN (SELECT b FROM S);
```



One loop, over  
the tuples of R

a	b
1	2
3	4

R

b	c
2	5
2	6

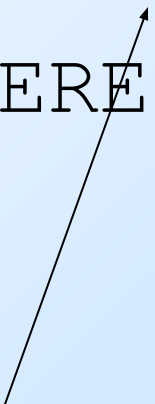
S

(1,2) satisfies  
the condition;  
1 is output once.

# This Query Pairs Tuples from R, S

```
SELECT a
FROM R, S
WHERE R.b = S.b;
```

Double loop, over  
the tuples of R and S



a	b
1	2
3	4

R

b	c
2	5
2	6

S

(1,2) with (2,5)  
and (1,2) with  
(2,6) both satisfy  
the condition;  
1 is output twice.

# The Exists Operator

- EXISTS(<subquery>) is true if and only if the subquery result is not empty.
- **Example:** From **Beers(name, manf)** , find those beers that are the unique beer by their manufacturer.

# Example: EXISTS

```
SELECT name  
FROM Beers b1  
WHERE NOT EXISTS (
```

Notice scope rule: manf refers to closest nested FROM with a relation having that attribute.

```
    SELECT *  
    FROM Beers  
    WHERE manf = b1.manf AND  
           name <> b1.name);
```

Set of beers with the same manf as b1, but not the same beer

Notice the SQL "not equals" operator

# The Operator ANY

- $x = \text{ANY}(\langle \text{subquery} \rangle)$  is a boolean condition that is true iff  $x$  equals at least one tuple in the subquery result.
  - $=$  could be any comparison operator.
- **Example:**  $x \geq \text{ANY}(\langle \text{subquery} \rangle)$  means  $x$  is not the uniquely smallest tuple produced by the subquery.
  - Note tuples must have one component only.



# The Operator ALL

- $x \text{ } \langle \rangle \text{ ALL}(\langle \text{subquery} \rangle)$  is true iff for every tuple  $t$  in the relation,  $x$  is not equal to  $t$ .
  - That is,  $x$  is not in the subquery result.
- $\langle \rangle$  can be any comparison operator.
- **Example:**  $x \geq \text{ALL}(\langle \text{subquery} \rangle)$  means there is no tuple larger than  $x$  in the subquery result.

# Example: ALL

- From **Sells(bar, beer, price)**, find the beer(s) sold for the highest price.

SELECT beer

FROM Sells

WHERE price  $\geq$  ALL(  


SELECT price

FROM Sells);

price from the outer  
Sells must not be  
less than any price.

# Union, Intersection, and Difference

- Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:
  - (`<subquery>`) UNION (`<subquery>`)
  - (`<subquery>`) INTERSECT (`<subquery>`)
  - (`<subquery>`) EXCEPT (`<subquery>`)

# Example: Intersection

- Using Likes(drinker, beer), Sells(bar, beer, price), and Frequents(drinker, bar), find the drinkers and beers such that:
  1. The drinker likes the beer, and
  2. The drinker frequents at least one bar that sells the beer.

Notice trick:  
subquery is  
really a stored  
table.

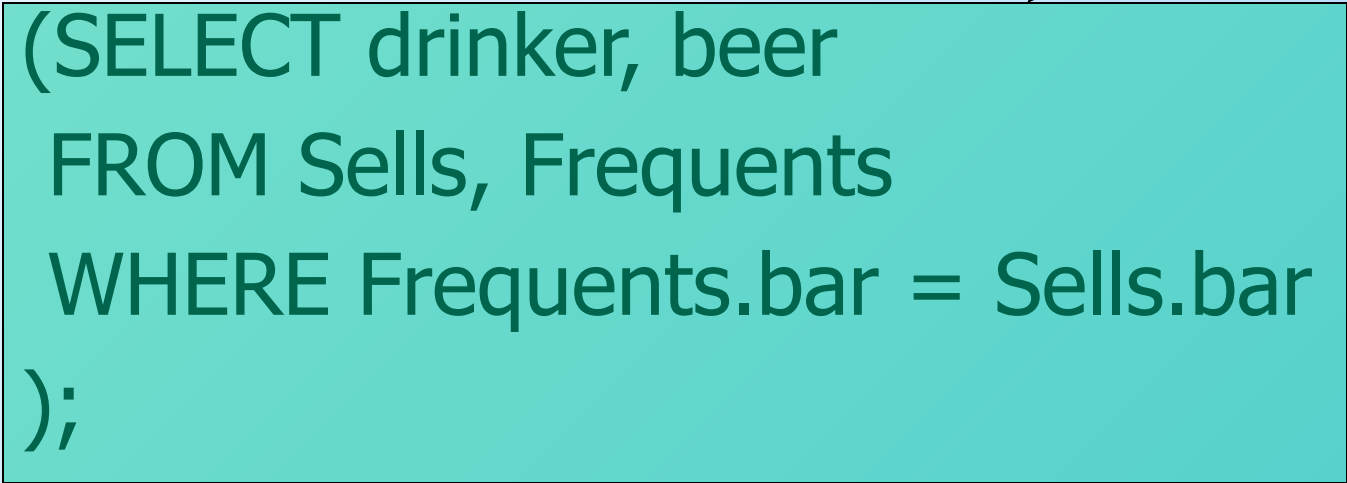
# Solution

The drinker frequents  
a bar that sells the  
beer.



```
(SELECT * FROM Likes)
```

INTERSECT



```
(SELECT drinker, beer  
FROM Sells, Frequents  
WHERE Frequents.bar = Sells.bar  
);
```

# Bag Semantics

- Although the SELECT-FROM-WHERE statement uses bag semantics, the default for union, intersection, and difference is set semantics.
  - That is, duplicates are eliminated as the operation is applied.

# Motivation: Efficiency

- When doing projection, it is easier to avoid eliminating duplicates.
  - Just work tuple-at-a-time.
- For intersection or difference, it is most efficient to sort the relations first.
  - At that point you may as well eliminate the duplicates anyway.

# Controlling Duplicate Elimination

- Force the result to be a set by  
`SELECT DISTINCT . . .`
- Force the result to be a bag (i.e., don't eliminate duplicates) by `ALL`, as in  
`. . UNION ALL . . .`



# Example: DISTINCT

- From `Sells(bar, beer, price)`, find all the different prices charged for beers:

```
SELECT DISTINCT price  
FROM Sells;
```

- Notice that without `DISTINCT`, each price would be listed as many times as there were bar/beer pairs at that price.

# Example: ALL

- Using relations **Frequents(drinker, bar)** and **Likes(drinker, beer)**:

```
(SELECT drinker FROM Frequents)
```

```
EXCEPT ALL
```

```
(SELECT drinker FROM Likes);
```

- Lists drinkers who frequent more bars than they like beers, and does so as many times as the difference of those counts.

# Join Expressions

- SQL provides several versions of (bag) joins.
- These expressions can be stand-alone queries or used in place of relations in a FROM clause.

# Products and Natural Joins

- Natural join:  
`R NATURAL JOIN S;`
- Product:  
`R CROSS JOIN S;`
- Example:  
`Likes NATURAL JOIN Sells;`
- Relations can be parenthesized subqueries, as well.

# Theta Join

- $R \text{ JOIN } S \text{ ON } \langle \text{condition} \rangle$
- **Example:** using  $\text{Drinkers}(\text{name}, \text{addr})$  and  $\text{Frequents}(\text{drinker}, \text{bar})$ :

```
Drinkers JOIN Frequents ON  
    name = drinker;
```

gives us all  $(d, a, d, b)$  quadruples such that drinker  $d$  lives at address  $a$  and frequents bar  $b$ .